

***MMN Queue Simulation  
Assignment Report***

**Distributed Computing**

Professor Matteo Dell'Amico

Seyede Aida Atarzadeh Hosseini - 6936800  
Delaram Doroudgarian - 5881909

Introduction .....	2
Crucial Parameters and Formula .....	2
Implementation of M/M/N Queue .....	2
Implementation of Supermarket Queueing .....	3
Theoretical VS Practical Plot .....	4
Implementation of M/M/N Queue with Weibull .....	5
Using Ring Topology in Supermarket to Choose Nodes Instead of Choosing Them Randomly ...	7
Optimizing Ring Topology by Adding Finger Table to Nodes .....	9
Comparison of the Plots Before and After Using the Ring Topology .....	10
Conclusion .....	12

## Introduction:

This assignment emphasizes the practical investigation of system behavior through simulation, exploring the complexities of distributed systems. It challenges conventional simplifying assumptions, such as memoryless exponential distributions and single-server models, by adopting a more nuanced approach. The tasks involve completing a discrete event simulation (DES) framework, implementing an M/M/1 FIFO simulation, and extending it to an M/M/n model. These efforts aim to provide practical experience in analyzing system behavior under realistic conditions. This report outlines the process, highlighting the methods employed, obstacles encountered, and key findings obtained throughout the exploration.

## Crucial parameters and formula:

**lambda( $\lambda$ ):** it represents the arrival rate of jobs into the system. It is typically measured as the number of arrivals per unit of time.

- Low  $\lambda$ : Jobs arrive infrequently, and queues are mostly empty.
- High  $\lambda$ : Jobs arrive quickly, and queues become congested.

**mu( $\mu$ ):** It represents the service rate, which is the rate at which jobs are processed or completed by a server. It is a key parameter that helps determine how quickly the system can handle incoming jobs. It's typically measured as the number of jobs completed per unit of time.

**Number of servers (n):** The total number of servers available in the system to process entities concurrently. The number of servers is a fundamental distinction between **M/M/1** and **M/M/N** queueing systems. In M/M/1 there is only one server. However, there is more than one servers in M/M/N.

**Maximum simulation time (max\_t):** This value ensures that the simulation does not run indefinitely and provides a boundary for processing events. max\_t limits the simulation to a finite duration and all events (arrivals, completions, and queue recordings) are processed up to this time.

**Times:** It records the lengths of all queues at specific intervals during the simulation. It serves as a log or history of queue states, used for analysis, such as plotting the queue length distribution.

**d:** d is the number of randomly chosen queues. When d = 1, it's a queue is randomly chosen from the available ones. For d > 1, a more advanced "Supermarket Model" is used. In this model, d queues are randomly selected, and the least busy queue is chosen for the job.

## Implementation of M/M/N Queue:

- Initially, the M/M/1 queue was implemented by completing the provided code framework.
- Subsequently, the implementation was extended to an M/M/n queue. To accommodate n servers, the sim.queue variable was modified from a single collections.deque to an array of collections.deque, corresponding to the number of servers. Additionally, the sim.running variable was updated from a single variable to an array to account for multiple servers processing jobs simultaneously.
- It is important to note that modifications were also made to the Arrival and Completion functions to support the changes in the queue and server structure.

```
self.running = [None] * n # if not None, the id of the running job (per queue)
self.queues = [collections.deque() for _ in range(n)] # FIFO queues of the system
```

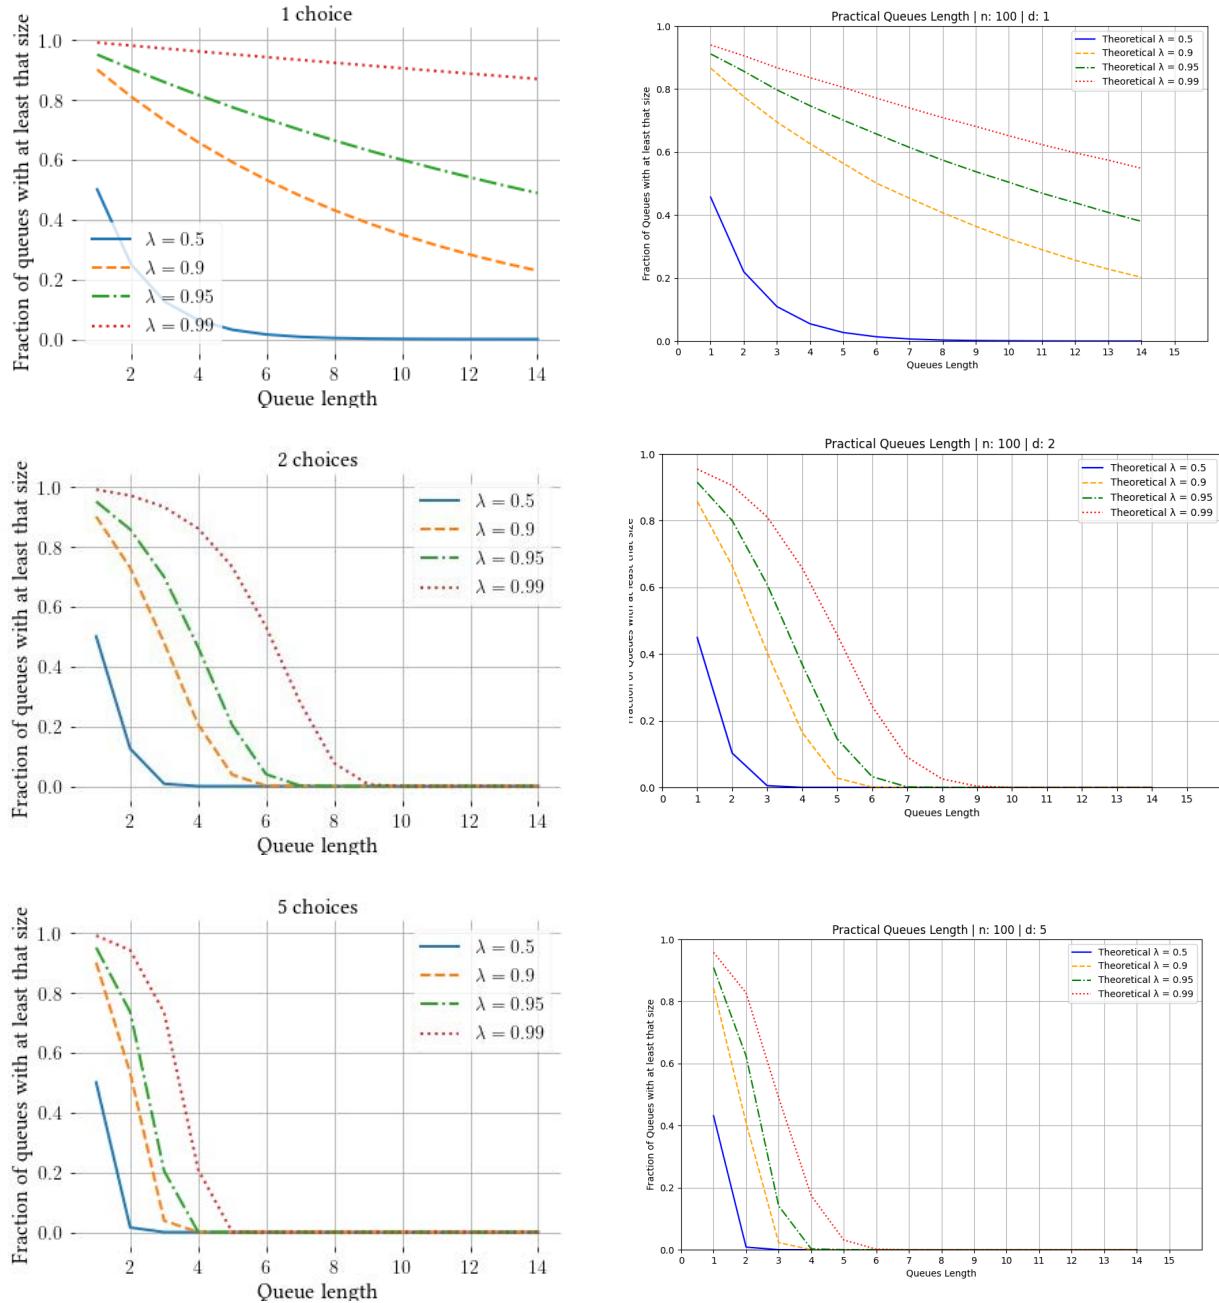
### Arrival Function with the Supermarket Model Optimization:

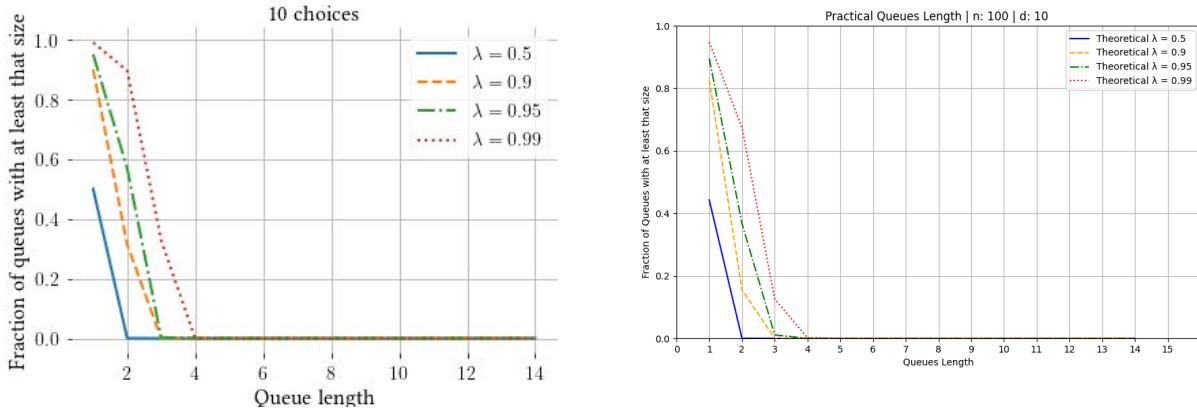
The Arrival function in the M/M/N queue simulation models the entry of new jobs into the system, ensuring optimal queue selection using the Supermarket Model optimization. When a job arrives.

```
class Arrival(Event):
    def __init__(self, job_id):
        self.id = job_id
    def process(self, sim: MMN):
        sim.arrivals[self.id] = sim.t
        samples = sample(sim.queues, sim.d)
        shortest_lists = [lst for lst in samples if len(lst) == min(len(sublist) for sublist in samples)]
        selected_list = choice(shortest_lists)
        queue_index = sim.queues.index(selected_list)
        if sim.running[queue_index] is None:
            sim.running[queue_index] = self.id
            sim.schedule_completion(self.id, queue_index)
        else:
            sim.queues[queue_index].append(self.id)
        sim.schedule_arrival(self.id + 1)
```

### Theoretical vs Practical plot:

We began by plotting both the theoretical plot (on the left) and the practical plot (on the right), using the parameters  $n=100$ ,  $\text{max\_t}=100$ . These parameters were consistently applied in all subsequent plots.





In advance of comparing the plots, let me explain the concept of “fraction of queues with at least  $X$  size”:

This concept refers to the proportion of queues that have a length greater than or equal to a specific size.

Fraction on queues represents the ratio percentage of queues that satisfy a certain condition (in this case, having a length greater than or equal to a specified value) relative to the total number of queues.

$$\text{Fraction} = \frac{\text{Number of queues with length } \geq X}{\text{Total number of queues}}$$

- This metric helps assess how well the system distributes the load across queues.
- If the fraction for longer queue lengths ( $\text{length } \geq 10$ ) is low, it indicates that there are only a few long queues and the load is distributed effectively.
- If the fraction for shorter queue lengths ( $\text{length } \geq 2$ ) is high, it shows that most queues are short, which suggests the system is lightly loaded or well-optimized.

## d=1

- When queues are chosen randomly (without considering their lengths), the system fails to distribute the load equally.
- At  $\lambda=0.9$ , the distribution is uneven, and there are more long queues compared to lower values of  $\lambda$ .
- For  $\lambda=0.5$ , the queues are generally shorter because the system experiences a lighter load.
- The practical model shows a similar trend but with a more uniform load distribution across the queues.

## d=2

- When two queues are randomly selected, and their lengths are compared, the load distribution improves significantly, resulting in fewer long queues.
- For higher values of  $\lambda$  ( $\lambda \geq 0.9$ ), long queues still exist but are noticeably less frequent compared to the  $d=1$  case.
- The practical outcomes closely match the theoretical predictions

**d=5**

- When d=5, the system distributes the load much better, significantly reducing the queue lengths.
- Long queues are almost nonexistent, even at high arrival rates ( $\lambda=0.99$ ).
- The practical model also shows shorter queues and improved load distribution with d=5.
- The simulated results are very similar to the theoretical predictions.

**d=10**

- With d=10, there is a slight improvement in load distribution compared to d=5.
- However, the difference is minimal because the system has already reached a balanced state.
- The practical model also shows that increasing d to 10 doesn't provide significant improvements over d=5.
- The practical and theoretical results are almost alike.

**Summary of observations:**

- Increasing d improves load balancing and reduces queue lengths.
- The most noticeable improvement is seen at d=5, and increasing d beyond this point results in only minimal changes.
- The simulation and theoretical plots generally match well.
- d=5 appears to be the optimal choice, providing a good balance between complexity and improved load distribution.

**Implementation of M/M/N queue with Weibull:**

we have implemented two different service time distributions: exponential and Weibull. The left-side plots represent the system behavior under the exponential distribution, while the right-side of plots demonstrate the system's performance when using the Weibull distribution. The comparison of these plots provides insights into the impact of different service time distributions on the system's behavior.

It is characterized by a **shape parameter** ( $k$ ) and a **scale parameter** ( $\lambda$ ).

- When  $k=1$ : The Weibull distribution becomes an **exponential distribution** (memoryless). An exponential distribution is a memoryless distribution, where the probability of an event occurring is independent of past occurrences.
- When  $k<1$ : It models a **heavy-tailed** distribution, meaning small values are common, but large values (outliers) occasionally occur.
- When  $k>1$ : It approaches a **bell-shaped** distribution, where values are concentrated around the mean.

The Weibull generator is defined in the workloads.py file, and it's used to produce interarrival times and service times in the simulator.

```
def weibull_generator(shape, mean):
    """Returns a callable that outputs random variables with a Weibull distribution having the given shape and mean."""
    return functools.partial(random.weibullvariate, mean / math.gamma(1 + 1 / shape), shape)
```

```
self.arrival_rate = lambd * n # frequency of new jobs is proportional to the number of queues
self.completion_rate = mu
self.arrival_gen=weibull_generator(weibull_shape,1/self.arrival_rate)
self.service_gen=weibull_generator(weibull_shape,1/self.completion_rate)
self.schedule(self.arrival_gen(), Arrival(0)) # schedule the first arrival
```

Arrival\_shape and service\_shape are shape parameters for the Weibull distribution.  
 $1/\text{self.lambd}$  and  $1/\text{self.mu}$  are the mean service times for the Weibull distribution.  
`self.arrival_gen()` is called whenever a new job needs to be scheduled for arrival, and  
`self_service_gen()` is called whenever a job's service time is required.

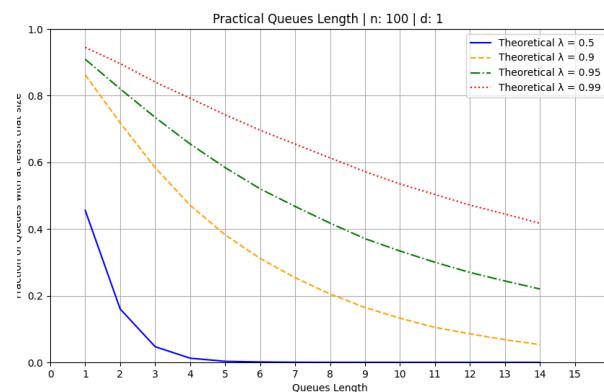
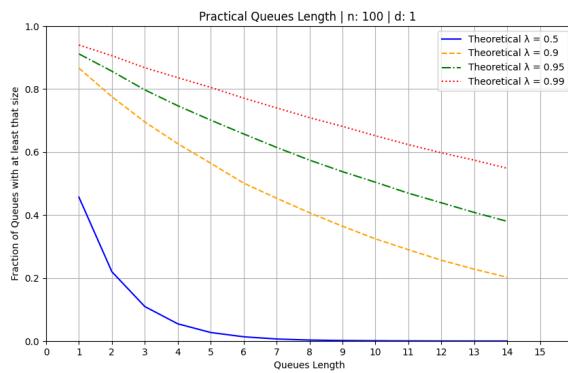
```
def schedule_arrival(self, job_id):
    self.schedule(self.arrival_gen(), Arrival(job_id))
```

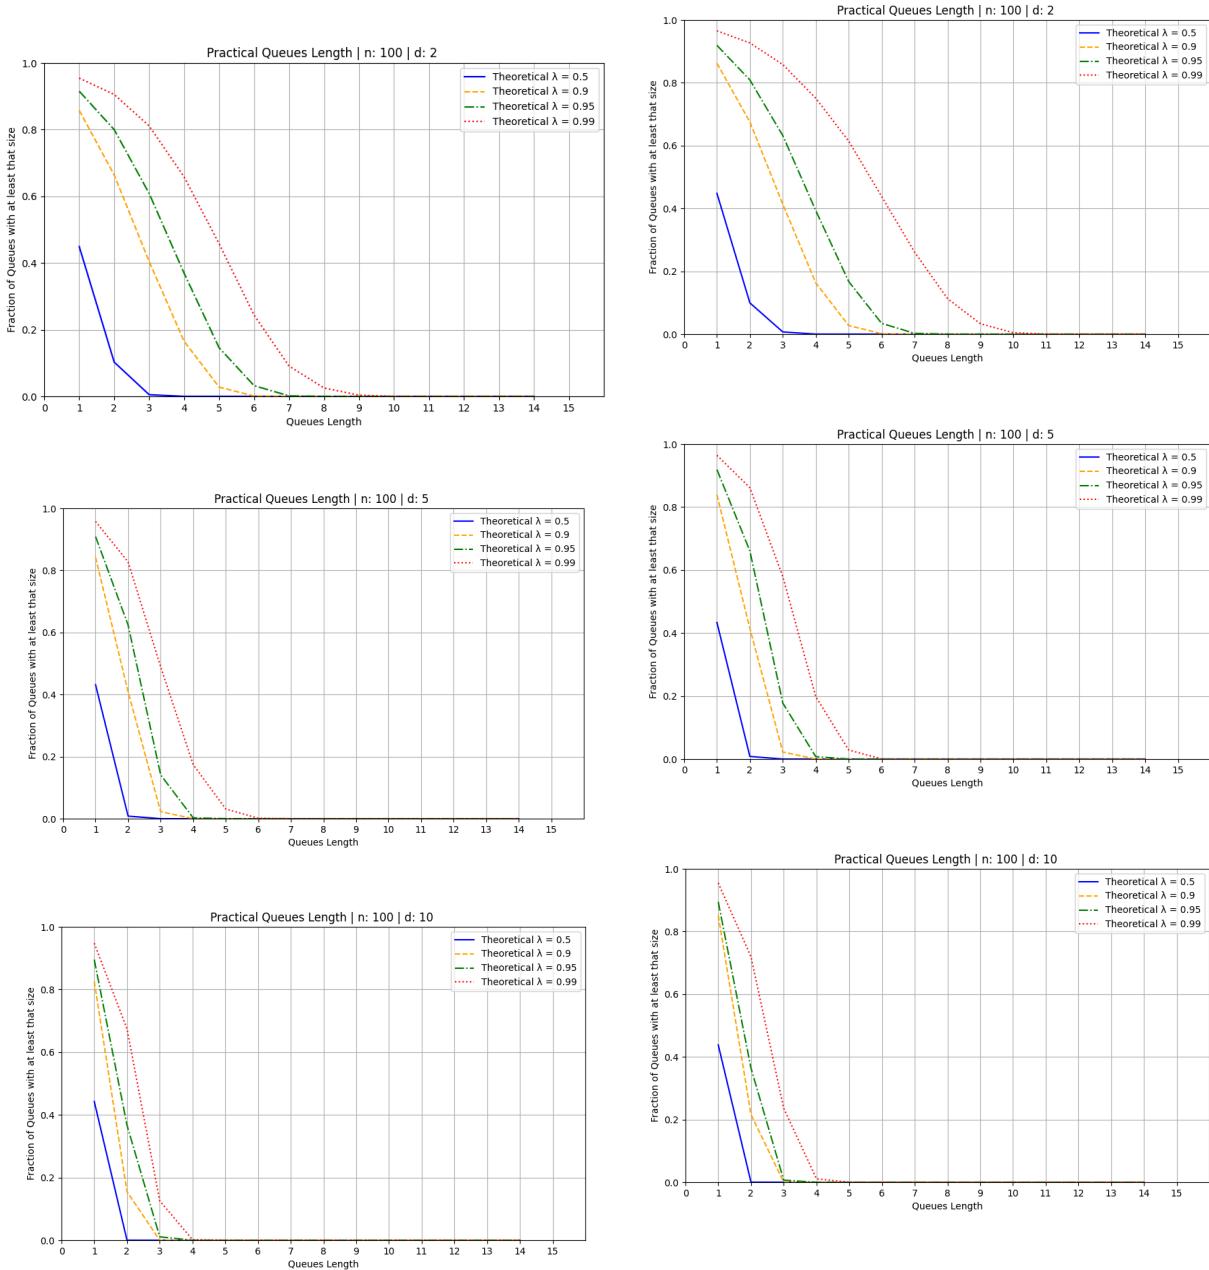
```
def schedule_completion(self, job_id, queue_index):
    self.schedule(self.service_gen(), completion(job_id,queue_index))
```

## Results analysis

The left-side plots represent the system's behavior before applying the Weibull distribution and the right-side ones are after.

### Shape = 1



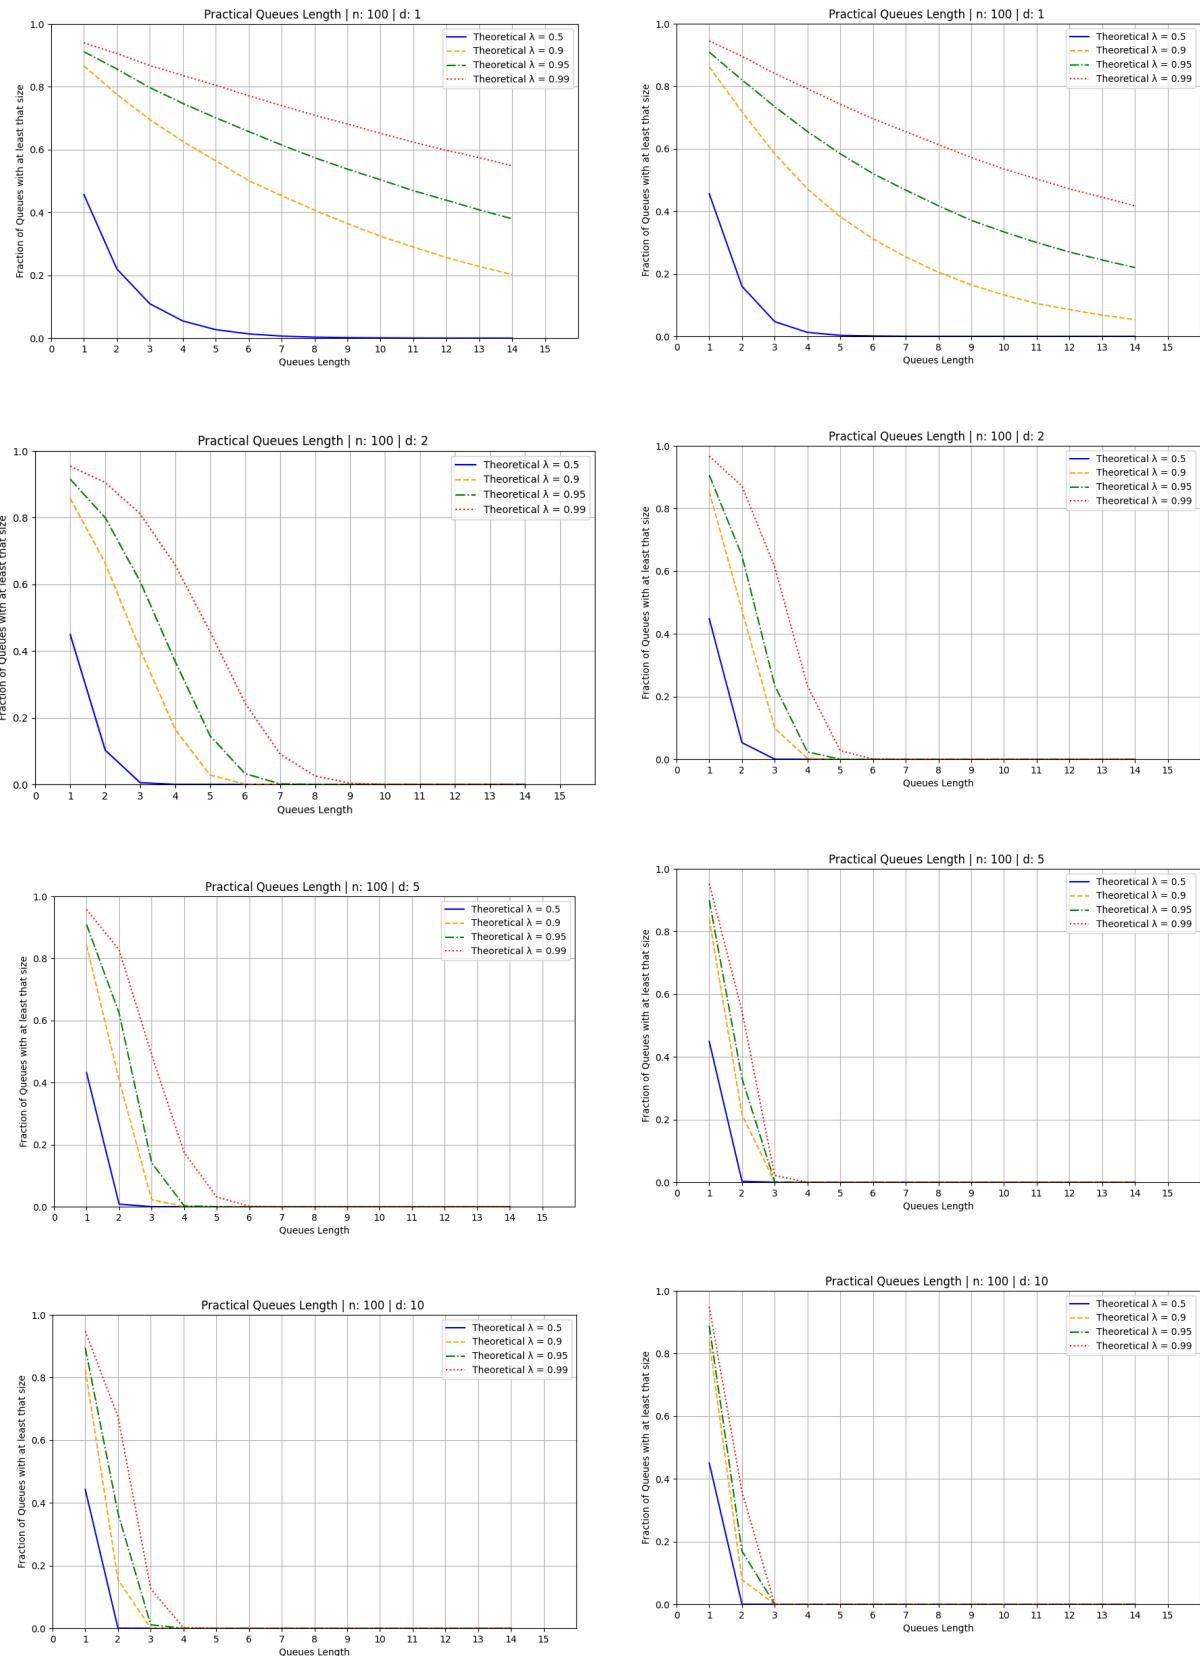


Since shape = 1 in the Weibull distribution corresponds to an exponential distribution, we expect similar behavior before and after applying Weibull.

The overall trend remains nearly the same between both sets of plots.

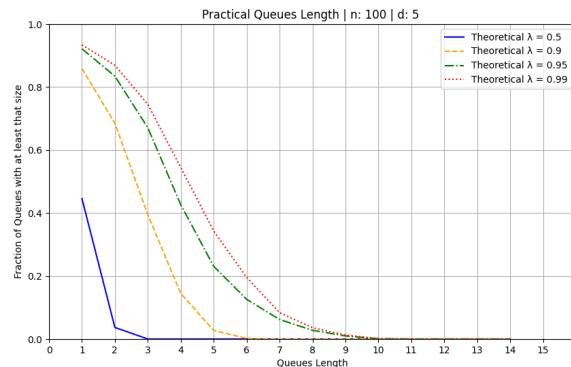
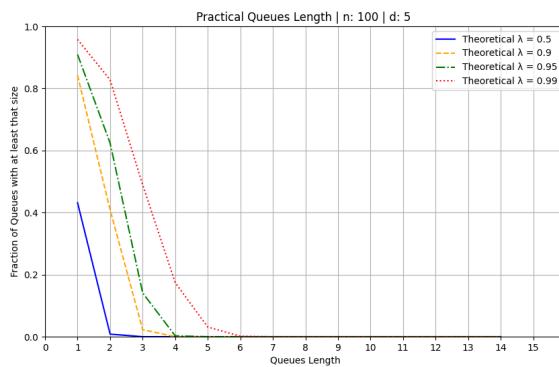
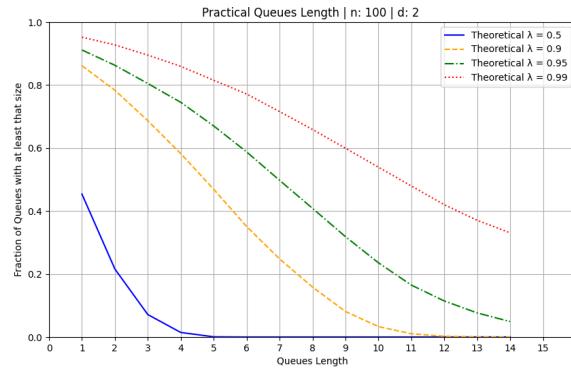
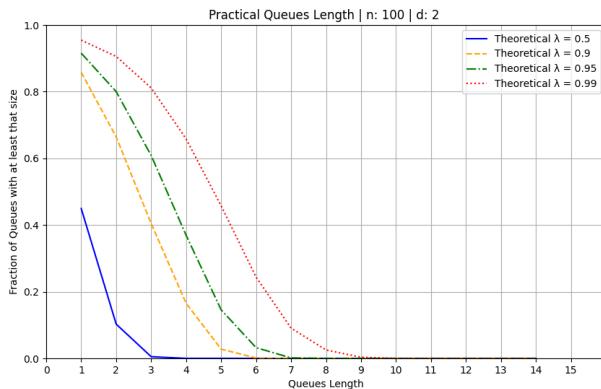
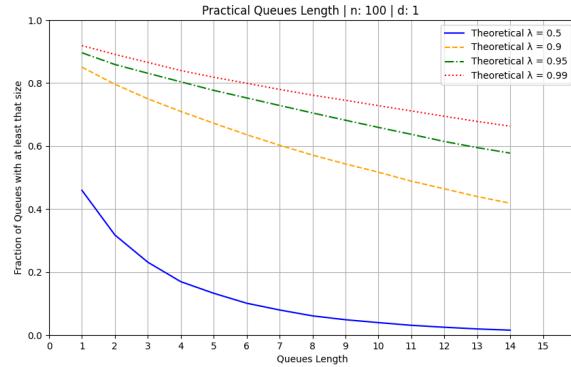
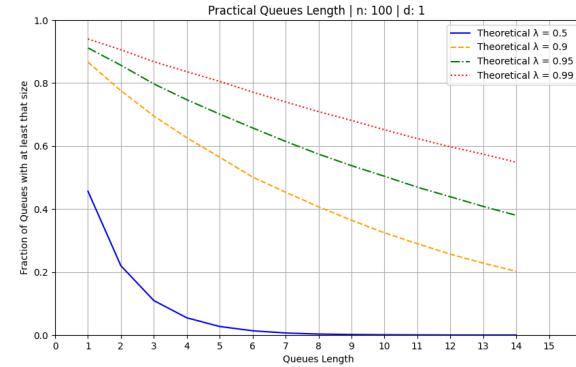
No shift is observed between the left and right plots, which confirms that Weibull (shape = 1) does not change queue behavior.

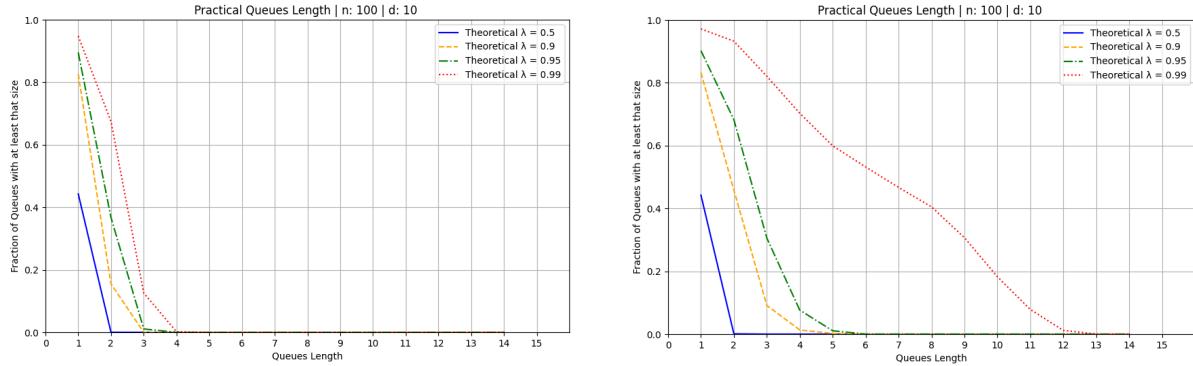
Shape > 1



The curves drop off more sharply, indicating that the number of long queues is reduced faster than in the exponential case. This aligns with the behavior of a bell-shaped Weibull distribution, where extreme values are less frequent.

Shape < 1





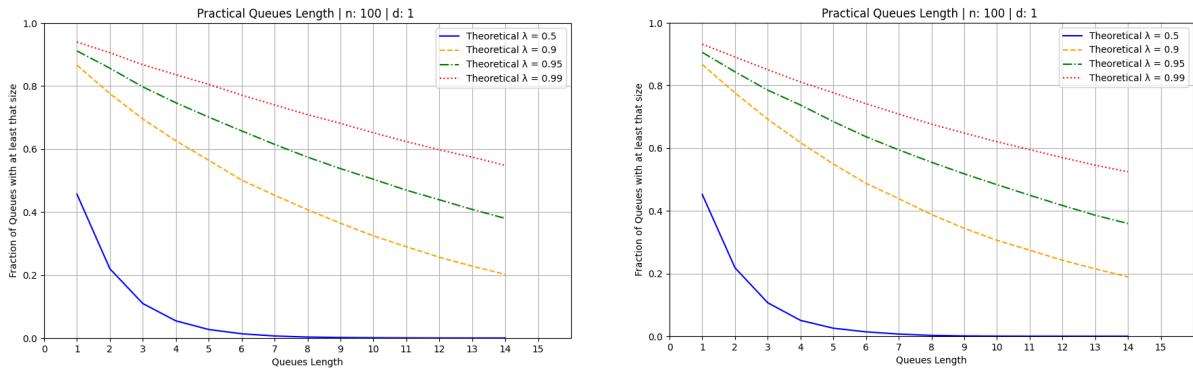
Using a Weibull distribution with shape  $< 1$  negatively affects queue performance by introducing extremely long jobs, which increase congestion and slow down processing.

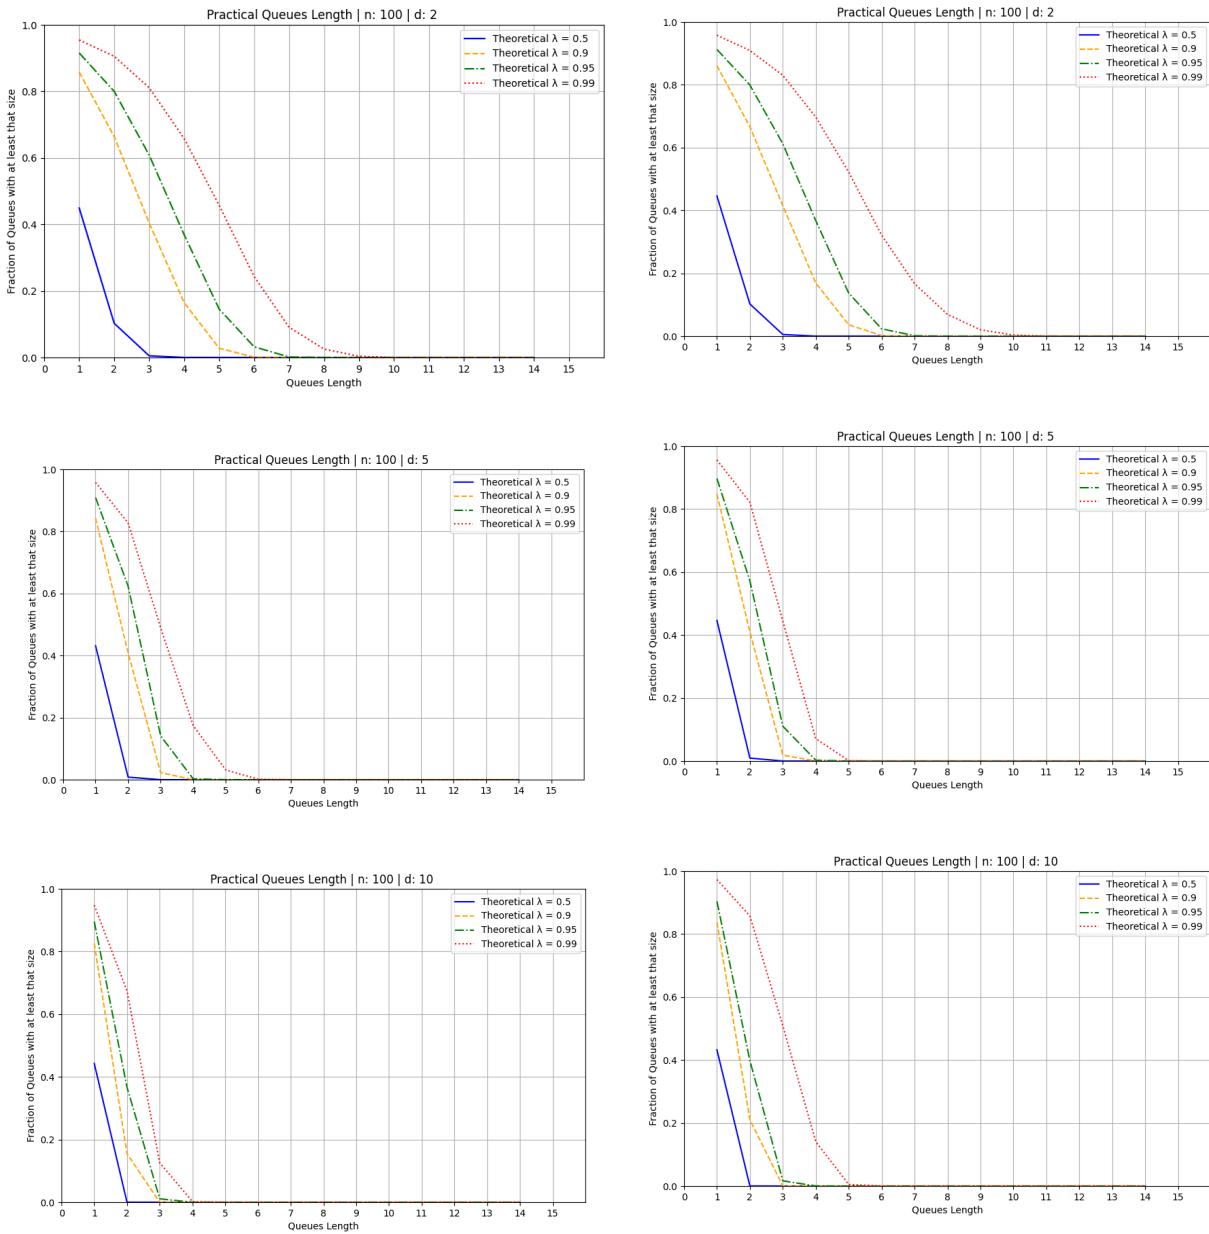
Even at high  $d$ , the system cannot fully compensate for the impact of long-tail jobs, meaning queues are more likely to grow.

The effect of increasing  $d$  is still beneficial, but it does not eliminate the inefficiency caused by heavy-tailed distributions.

### Job's Priority Extention:

The current implementation lacks support for job priorities, preventing the prioritization of more urgent tasks. The goal is to ensure high-priority jobs are placed at the top of the heap queue. To evaluate the impact of the modification, we compared the results from running the simulator both without and with the priority concept.





**d=1:**

- Prioritization cannot optimize job assignments effectively.

**d=2:**

- Still not very effective, as the system does not have enough flexibility to fully benefit from prioritization.

**d=5:**

- Now prioritization shows a clear improvement.
- High-priority jobs get processed faster, reducing their waiting times.
- This is the best balance between prioritization and queue selection.

d=10:

- Prioritization impact decreases.
- Since queue selection is already optimal, prioritization doesn't provide much additional benefit.
- The system's natural efficiency with high d overcomes the need for prioritization.

Final Consideration:

- Prioritization is most beneficial at d=5, where it effectively reduces waiting times for high-priority jobs while still benefiting from queue selection.
- For very low d (d=1,2), prioritization has little effect because jobs accumulate in queues with minimal optimization.
- For very high d (d=10), prioritization becomes unnecessary since the queue selection process is already highly efficient.
- Best configuration: Combining prioritization with d=5 for optimal performance!